
AxoSim: Learned Neuron Models for Fast Population Simulation and Inference-Time Adaptation

Davide Wiest

Axym Labs, Darmstadt, Germany

August 2026

Code: [repository](#)

Project page: [AxoSim](#)

Correspondence: contact@axym.org

Abstract

Detailed neuron simulators resolve morphology and ion-channel mechanisms, but repeated numerical integration is costly at population scale. Learned surrogates reduce this cost; in recurrent circuits, however, each prediction also affects the model’s future inputs. AxoSim combines morphology-conditioned GRU and Mamba neuron models with delayed synaptic routing, mutable efficacies, and differentiable population state. Under the standard AxoBench evaluation, AxoSim-GRU Medium reaches spike Mean F1 0.522, compared with up to 0.077 for released Branch-ELM models, and reduces Voltage and Dynamics SERA errors by 21.3% and 4.4%, while processing 16.8 million neuron-steps per second in the matched inference benchmark. On an RTX 5090, the runtime simulates 65,536 neurons with 4,000 retained contacts per neuron at $3.38\times$ real-time speed. On the same consumer-grade GPU, populations of 20,000 neurons are trained in real time with full backpropagation through time (BPTT) at 2,000 retained contacts per neuron. In a temporal-shift experiment on Spiking Heidelberg Digits (SHD), morphology and synaptic adaptation raise target accuracy by 6.46 ± 1.20 percentage points across six independently trained models. A separate corpus provides 50,000 validated population traces from 100 recurrent connectomes with connectome-disjoint splits. Together, these experiments separate teacher-forced fidelity from recurrent simulation and inference-time adaptation.




Figure 1: Five quantities available for all three systems. AxoSim shows Mamba Medium and the throughput-leading GRU Small; CoreNEURON’s self-comparison defines perfect fidelity. Raw values, mappings, and run-time contracts appear in Appendix A.5.

Table of contents

1	Overview	3
2	System design	5
3	Evaluation methodology	8
4	Single-neuron fidelity and efficiency	9
5	Connected population simulation	11
6	Inference-time adaptation	14
7	Relation to learned and mechanistic simulators	17
8	Discussion and future work	18
9	Conclusion	20
A	Measurement contracts and supplementary evidence	22

1 Overview

1.1 The simulation problem

Biophysically detailed neuron simulators integrate membrane mechanisms over a branched morphology, which preserves a direct relation between model parameters and cellular processes. That detail is indispensable when the question concerns a named conductance, compartment, or synaptic mechanism, but it also makes repeated simulation and gradient-based fitting expensive as the number of cells grows. Learned neuron models instead fit responses generated by a detailed teacher and approximate the map from synaptic input histories to somatic voltage and spikes [2, 15].

Most evaluations of learned neurons hold the input trace fixed while comparing the predicted and teacher responses. This teacher-forced evaluation measures the quality of a single response. A recurrent population introduces an additional difficulty because a predicted spike is delayed, delivered to downstream dendrites, and incorporated into the input for later predictions. An error can therefore change the current output and the distribution of future inputs. Population simulation must measure this feedback directly because single-neuron accuracy omits its recurrent consequences.

AxoSim is a learned-neuron simulation platform designed around that distinction. It complements the mechanistic teacher with compatible interfaces for fitting neuron responses, connecting the fitted models through explicit synaptic contacts, advancing persistent populations, and differentiating through a supplied trajectory. Its learned coordinates are predictive variables rather than identified ion-channel mechanisms. CoreNEURON simulations of Hay layer-5 pyramidal cells provide the principal detailed-neuron targets, while released Branch-ELM models provide the principal learned reference.

1.2 From single neurons to recurrent populations

At the single-neuron level, each AxoSim model receives events that preserve their sign and dendritic destination, combines them through a morphology-conditioned front end, and predicts a somatic-voltage coordinate and spike logit at every native millisecond. GRU and Mamba temporal cores share this input and output contract, so their fidelity and execution cost can be compared directly. The full models prioritize response fidelity. The compact AxoSim-Lite specialization preserves the same event semantics while reducing state and computation for persistent connected populations.

The population runtime stores recurrent neuron state and a directed graph of synaptic contacts. Each contact preserves its source, sign, branch destination, delay, and mutable efficacy. During each simulation step, thresholded model outputs define spikes. The runtime routes those events after their delays, accumulates their efficacies at the target branches, and records requested telemetry. PyTorch autograd differentiates through the same operations. For large populations, the fused GPU implementation changes their scheduling and memory layout.

Parameter ownership determines the scale of adaptation. Shared model weights define the learned response function; behavior parameters belong to individual neurons, morphology parameters are shared by a structural class, and synaptic parameters belong to individual contacts. Any selected group can receive gradients through the complete supplied sequence. This separation allows the same interface to express morphology transfer, changes in synaptic efficacy, and task adaptation while preserving their interpretation as predictive coordinates.

1.3 Evidence across fidelity, execution, and adaptation

AxoBench introduced spike Mean F1 across 0–5 ms matching tolerances, Voltage SERA, and Dynamics SERA as the core metric set for single-neuron evaluation [21]. Figure 2 relates these three quantities to parameter count, forward-pass computation, and inference throughput. The largest separation from Branch-ELM occurs in spike timing. Voltage SERA also decreases substantially, whereas the smaller Dynamics-SERA change identifies native-time voltage increments as the harder quantity.




Figure 2: Accuracy–cost tradeoffs for AxoSim and released Branch-ELM models. Each fidelity metric—Voltage SERA, Dynamics SERA, and spike Mean F1—is evaluated as a function of trainable parameters, forward-pass FLOPs, and inference throughput. Lines connect models from the same architecture family, and orange arrows link the Pareto-selected Branch-ELM and AxoSim models. FLOPs are the operations recognized by `torch.profiler` for a batch-one, 500-step forward pass with 1,278 input channels. Throughput is measured on an RTX 5090 with batch 32, 500 steps, and float32. The Pareto rule is defined in Appendix A.6.

The frontier contains three distinct leaders. GRU Medium records the highest Mean F1 and lowest Voltage SERA, GRU Small provides the highest inference rate, and Mamba Small records the lowest Dynamics SERA. Mamba Medium nearly matches the spike leader at greater state and computational cost. These values measure responses to held-out teacher inputs. Population timing covers a complete recurrent step at explicit neuron and contact counts, and connected diagnostics measure activity after predictions enter the feedback loop. Finally, adaptation analyses test transfer to an excluded morphology, a controlled conductance shift, and a temporal distribution shift on SHD. An internal-state analysis additionally tests whether controlled interventions are organized by their detailed-teacher response. Table 1 associates each result with its measurement condition.

Table 1: Principal AxoSim results and the conditions attached to each measurement.

Question	Result	Measured condition
Neuron fidelity	Mean F1 0.522; the lowest AxoSim Voltage SERA is 21.3% below its Branch-ELM reference, and the lowest Dynamics SERA is 4.4% below its reference.	Ordinary AxoBench, held-out morphology-balanced traces.
Intervention geometry	Temporal-state geodesics correlate 0.446 with detailed-teacher response distance, compared with 0.211 before temporal integration.	12,500 paired traces across five interventions and five morphologies.
Population execution	1,048,576 neurons advance at 1.177 billion neuron-steps/s with 362 incoming synaptic contacts per neuron; 65,536 neurons with 4,000 contacts advance at 3.38 times real time.	Complete RTX 5090 step with delayed routing, mutable efficacies, state, and 1% telemetry.
Population data	50,000 detailed-teacher traces from 100 recurrent connectomes assigned to disjoint context splits.	Five excitatory morphologies and one inhibitory morphology; complete connectomes define the split.
Adaptation	SHD target accuracy improves by 6.46 ± 1.20 percentage points; 20,000 neurons with 2,000 retained contacts complete a 256-step full-BPTT update in real time.	Six independently trained SHD source models; complete full-BPTT update on one RTX 5090.
Biological transfer	A 58-value adapter improves spike and dynamics fidelity on five excluded morphologies; synaptic parameters recover one controlled conductance shift.	Predictive transfer is established; cellular identification requires further controlled interventions.

AxoSim’s full neuron models improve held-out single-neuron prediction, especially spike timing. The connected runtime separately advances 1,048,576-neuron populations at the stated connectivity and differentiates through smaller, more densely connected populations. Because a prediction routed through a circuit can alter subsequent activity, recurrent validity requires connected measurements in addition to single-neuron fidelity. The evaluation consequently separates response fidelity, connected execution, and adaptation.

2 System design

2.1 A contact-resolved learned neuron

At native millisecond t , a neuron receives synaptic events $x_{t,k}$. Each event carries its sign, dendritic destination, delay, and positive efficacy. For morphology descriptor m , neuron-specific behavior coordinates a , and recurrent state h_t , AxoSim evaluates

$$(\hat{\ell}_t, \hat{v}_t, h_{t+1}) = f_\theta(x_{1:t}; m, a, h_t),$$

where $\hat{\ell}_t$ is the spike logit and $\hat{v}_t$ is the somatic-voltage coordinate defined by the data contract. For the standard NeuronIO conversion, $\hat{v}_t = (v_t + 67.7)0.1$, and evaluation converts the coordinate back to millivolts. A sigmoid converts the logit to a probability when needed, and thresholding the logit produces the discrete event used by the population runtime. Both outputs are defined at the native time step, which makes voltage trajectories and spike times available directly at native resolution.




Figure 3: Neuron model architectures and a measured response. (a,b) Computation proceeds upward from morphology-conditioned branch aggregation to residual temporal updates and separate output heads. Branch routing and aggregation follow Branch-ELM. Mamba Medium repeats the illustrated block four times before causal temporal refinement; γ scales residual updates, and c_v couples voltage into the spike logit. (c) Dashed lines associate the Mamba stages with one held-out 500-ms trace. Outputs and temporal PC 1 are shown at 1-ms resolution; input events use 10-ms counts. Gray ticks mark detailed-simulator spikes.

The branching front end in Figure 3 follows the central Branch-ELM abstraction [15]. In both model families, synaptic channels are assigned to dendritic branches, contributions are integrated within each branch, and the resulting branch representation supplies the temporal neuron state. AxoSim augments this shared structure with morphology-conditioned synaptic gains, a nonlinear learned branch projection, separate spike and voltage heads, and, in the Mamba family, causal temporal refinement. Comparisons with released Branch-ELM models therefore retain the common branching mechanism while measuring the combined value of the changed parameterization and temporal model.

After branch projection, gated recurrent updates provide temporal integration in AxoSim-GRU, while AxoSim-Mamba uses selective state-space blocks [6]. A GRU applies one state transition after another. The selective state-space recurrence additionally admits a parallel scan, which processes a length- T

history with $O(T)$ work and $O(\log T)$ parallel depth, whereas its streaming form retains constant-size state and requires $O(1)$ work for each new timestep. Complete-sequence evaluation uses the parallel form, and streaming inference uses the recurrent form. Both model families place their temporal transformation on a pre-normalized residual path because recurrent optimization remains sensitive to vanishing and exploding gradients [11],

$$x_{\ell+1} = x_{\ell} + \gamma_{\ell} \mathcal{M}_{\ell}(\text{LN}(x_{\ell})),$$

where $\mathcal{M}_{\ell}$ denotes the GRU or selective state-space transformation. The identity path spans the model depth, and a learned channel-wise multiplier γ_{ℓ} controls the residual update.

2.2 Persistent populations with delayed contacts

The compact P4 recurrence emits four distinct native-time predictions from one state update,

$$(\hat{y}_{4b}, \hat{y}_{4b+1}, \hat{y}_{4b+2}, \hat{y}_{4b+3}) = F_{\theta}(x_{\leq 4b-1}, h_b).$$

The four outputs are used directly. Population contacts have a minimum delay of four milliseconds, so each predicted event is available before it can affect another neuron. This causal relation permits blockwise execution while preserving the native-time event semantics.

A population combines persistent neuron states with a directed, delayed contact graph. Every synaptic contact stores its source neuron, sign, destination branch, delay, and mutable efficacy. A P4 advance consumes delayed branch inputs, updates neuron state, compacts emitted events, routes them through the graph, and accumulates their efficacies at destination branches. Requested trajectory samples move asynchronously to the host. The learned spike output and its threshold determine firing.

The PyTorch layer implements these operations with ordinary tensor programs and autograd. AxoRuntime uses phase- and tile-major layouts, an event ring, exact modulo delay routing, and fused compact-neuron kernels. Their scheduling differs, but both layers share the same contact and state definitions.

2.3 Parameter ownership and inference-time adaptation

Shared neuron-model weights are distinct from three adaptive parameter groups. Behavior coordinates belong to individual neurons and alter their encoded input sensitivity, memory timescale, recurrent proposal, and output calibration. Morphology coordinates are shared by neurons assigned to the same structural class. Synaptic coordinates belong to individual contacts and set positive efficacies before branch accumulation. These ownership rules allow a population to specialize at neuron, class, or edge scale while the shared model remains fixed.

Table 2: Ownership and role of the adaptive parameter groups in the compact population model.

Group	Ownership	Compact parameters	Effect
Behavior	one neuron	88 per neuron	temporal response and readout
Morphology	one structural class	88 per class	class-shared response
Synaptic	one synaptic contact	one per contact	positive input efficacy

Any subset of these groups participates in ordinary PyTorch autograd over the complete supplied sequence. The accelerated path captures input processing, forward execution, loss evaluation, full-sequence backward propagation, gradient clipping, and the optimizer step in one CUDA Graph. Each replay receives new events and targets and updates the selected parameter banks. Graph replay fixes

tensor dimensions and memory addresses while supporting repeated updates with fresh data at that shape.

Sparse event indices and values replace a dense neuron–time–contact tensor. When full-sequence differentiation would otherwise exceed memory, neuron chunks share a morphology adapter and contribute neuron-count-weighted losses before one optimizer step. CPU tests match the behavior, morphology, and synaptic gradients of chunked and unchunked execution.

3 Evaluation methodology

3.1 Separating fitting, selection, and evaluation

Each comparison separates data used to fit parameters, choose a model or adaptation schedule, and estimate final performance. The unit of separation follows the scientific question. Ordinary AxoBench uses morphology-balanced traces, population specialization uses recurrent contexts, SHD uses disjoint examples and independently trained source models, and connected evaluation uses unseen connectomes. Appendix A.1 records the exact partitions.

For the ordinary comparison, development data fix the model identity and spike threshold. The first 24 final traces determine one prevalence-matched threshold, and the remaining 96 provide all reported fidelity metrics. Released Branch-ELM models are evaluated unchanged.

Population-input specialization selects one learning-rate schedule per architecture family on development contexts and applies it unchanged on the final contexts. SHD selects the adaptive groups, learning rate, and update count on one source model, then reuses that choice for five independently trained sources. Connected-population comparisons pair alternatives on each unseen connectome. Final data are therefore used only after model and schedule choices have been fixed.

3.2 Neuron and population fidelity

The single-neuron evaluation follows the AxoBench core metric set [21]. Spike Mean F1 averages event F1 over integer matching tolerances from 0 through 5 ms,

$$\text{MeanF1}_{0:5} = \frac{1}{6} \sum_{\tau=0}^5 F_1(\tau).$$

Predicted and teacher events are matched one-to-one in temporal order within $\pm\tau$. The six constituent values distinguish exact event recovery from agreement that appears only at a permissive tolerance.

For prediction error $e_i = \hat{y}_i - y_i$ and teacher-derived relevance $r_i \in [0, 1]$, the conditional squared error above threshold u is

$$E(u) = \frac{1}{|\{i : r_i \geq u\}|} \sum_{i:r_i \geq u} e_i^2.$$

SERA is the support-normalized integral of $E(u)$, evaluated at 101 evenly spaced relevance thresholds. Voltage SERA applies the construction to somatic voltage and has units of mV^2 . Dynamics SERA applies it to adjacent native-time voltage changes and has units of mV^2/ms^2 . SERA already names the squared quantity, and Root-SERA denotes its square root. Lower SERA is better.

Connected evaluation keeps three discrepancies separate. A morphology term compares morphology-conditioned firing rates with fixed targets, a temporal term compares input-count autocorrelations at lags from 1 to 256 ms, and a stationarity term measures full-to-tail drift in global and role-conditional activity. Global rate, input load, variability, and active fraction are also checked for collapse or composition changes. The equations and fixed normalizations appear in Appendix A.10.

3.3 Uncertainty and runtime contracts

Single-neuron differences use 1,000 paired bootstrap resamples of the 96 scored traces, stratified by morphology. All three metrics are recomputed from each paired sample. These intervals describe trace-level uncertainty for fixed trained models. Retraining variation requires independently fitted models, as used for adaptation, where the report gives the mean and sample standard deviation across source models. Population-training results use the paired outcome on every unseen connectome. The reported uncertainty therefore follows the source of variation in each analysis.

Runtime quantities are attached to explicit workloads. Isolated-neuron rates use float32 and matched batch and sequence shapes. Population inference uses the p90 latency of a complete step, including event discovery, delayed routing, branch accumulation, mutable efficacy, recurrent state, and 1% telemetry. Population adaptation measures a complete 256-step update, including transfer, sparse input construction, forward execution, loss, full-sequence backward propagation, clipping, and optimization. Hardware, precision, repeats, and exclusions are tabulated in Appendix A.2.

CoreNEURON, Branch-ELM, and AxoSim use different workloads and hardware, so Figure 1 compares their directly measured quantities. A cross-system speed ratio is reserved for a future shared-workload measurement. Raw values and axis mappings appear in Appendix A.5.

4 Single-neuron fidelity and efficiency

4.1 Held-out fidelity and the practical model frontier

The largest quantitative separation in the ordinary AxoBench results occurs in spike timing (Table 3). AxoSim-GRU Medium reaches Mean F1 0.522, whereas the strongest released Branch-ELM spike reference reaches 0.077, an absolute difference of 0.445. At the same time, GRU Medium lowers Voltage SERA from 307.18 to 241.61 mV², a 21.3% reduction. Mamba Small lowers Dynamics SERA from 448.57 to 428.72 mV²/ms², a smaller 4.4% reduction. The dominant improvement is therefore event recovery, accompanied by a substantial voltage-trajectory gain and a modest improvement in native-time voltage changes.

Table 3: Held-out ordinary-AxoBench results. Lower SERA and higher Mean F1 and throughput are better. Inference is measured in millions of neuron-steps/s.

Model	Parameters	Voltage SERA	Dynamics SERA	Mean F1	Inference
AxoSim-GRU Small	17,432	243.92	429.12	0.486	21.79
AxoSim-GRU Medium	122,744	241.61	429.13	0.522	16.81
AxoSim-Mamba Small	32,536	244.29	428.72	0.454	18.86
AxoSim-Mamba Medium	1,128,200	244.46	439.00	0.520	4.40
Branch-ELM M30	10,952	307.41	448.57	0.077	0.742
Branch-ELM M50	19,252	307.18	448.58	0.076	0.735

For AxoSim-Mamba Medium, Mean F1 is 0.443 higher than the strongest Branch-ELM spike reference, with a paired morphology-stratified 95% bootstrap interval of [0.422, 0.463]. Its Voltage-SERA difference from the strongest released voltage reference is -62.72 mV^2 , with a 95% interval of $[-69.06, -56.72]$. AxoSim-GRU Medium has the highest final Mean-F1 point estimate, 0.0016 above Mamba Medium. The near tie in spike fidelity is practically important because the GRU model is much smaller and nearly four times faster at inference.

The target quantity determines model choice (Figure 2). GRU Medium leads Mean F1 and Voltage SERA, GRU Small leads inference throughput, and Mamba Small records the lowest Dynamics SERA. Mamba Medium provides a state-space alternative within 0.002 of the spike leader, at 4.40 million neuron-steps/s. The three leaders define the measured frontier. The larger Mamba state offers temporal capacity at a higher execution cost.

4.2 Where the fidelity gains arise

Mean F1 averages six spike-matching tolerances, and AxoSim’s advantage appears at every setting. GRU Medium exceeds Branch-ELM M30 at every tolerance in Figure 4. At exact timing, F1 is 0.261 versus 0.047, and at 5 ms it is 0.641 versus 0.102. The throughput-leading GRU Small has the same qualitative separation from the closest-size Branch-ELM reference.

Conditional squared-error curves provide the complementary voltage result. Integration of each conditional squared-error curve yields SERA. Across most of the relevance support, AxoSim has lower conditional voltage error than Branch-ELM, so the Voltage-SERA reduction is distributed across a broad threshold range. The smaller separation between the dynamics curves agrees with the modest aggregate Dynamics-SERA improvement. Thus, event recovery can improve markedly even when error in native-time voltage changes decreases only slightly.




Figure 4: Constituents of the core fidelity metrics. Spike F1 is reported across matching tolerances, followed by the conditional squared-error curves integrated by Voltage and Dynamics SERA.

4.3 Internal representations of controlled interventions

Paired interventions permit a direct test of the information retained inside the temporal state. We evaluated AxoSim–Mamba Medium on all 12,500 paired traces in axobench-interventions. Each point in Figure 5a averages a deterministic, disjoint mini-cohort of seven or eight pairs from one intervention–morphology shard. The resulting signature concatenates the mean state displacement in eleven successive 500-ms windows after the initial 500 ms. A three-dimensional UMAP is used for static visualization [10]. All quantitative comparisons use the original representation space.

The 1,575 cohort signatures form compact, intervention-dependent regions, with marker shape preserving morphology identity. A separate linear coordinate system is used for time in Figure 5b. Its first three principal components explain 83.2% of the state-displacement variance, and the colored paths show condition means at 500-ms intervals. Gray arrows average the next-window displacement of individual pairs within occupied spatial cells, following established low-dimensional analyses of recurrent dynamics [3, 17]. Across the 25 intervention–morphology groups, within-intervention di-

rections exceed between-intervention directions by 0.563 in cosine similarity, with a shard-bootstrap 95% interval of [0.559, 0.566].

Biological interpretation requires an external response quantity in addition to visual separation. Following the external-ruler construction used for scientific-model manifolds [12] and representational similarity analysis [8], angular distances define the smallest connected nearest-neighbor graph in the original state space. The shortest-path distances correlate 0.446 with detailed-teacher voltage-and-spike response distance across 300 group pairs, with a 95% interval of [0.443, 0.451]. Before temporal integration, the branch representation reaches 0.211 [0.167, 0.213], while the input-perturbation control reaches 0.471 [0.443, 0.498]. Temporal integration therefore reorganizes the branch features toward the causal response geometry and approaches the 0.471 association already present in the intervention input.




Figure 5: Internal organization of controlled AxoBench interventions. (a) Three-dimensional UMAP of 1,575 disjoint mini-cohort state signatures; color denotes intervention and marker shape denotes morphology. (b) Separate PCA trajectories at 500-ms intervals. Small gray arrows show the average next-window motion within occupied cells of a $4 \times 4 \times 4$ state-space grid. (c) Group-level manifold geodesics versus the detailed teacher’s voltage-and-spike response distance. The 300 points comprise every pair of 25 intervention–morphology groups. Quantitative distances use the original representations. The three-dimensional coordinates serve visualization.

5 Connected population simulation

5.1 Throughput as a function of connectivity

The population benchmark advances 1,048,576 heterogeneous neurons at native 1 ms resolution. One real-time second therefore requires 1.048576×10^9 neuron-steps. The measured step includes event-derived firing, delayed contact routing, branch accumulation, mutable efficacies, recurrent state, and pinned-host telemetry for 1% of the neurons.

Throughput decreases as incoming synaptic connectivity increases (Figure 6). The million-neuron measurement uses 362 incoming contacts per neuron and advances at 1.177×10^9 neuron-steps/s. At 4,000 contacts, a closed-loop 65,536-neuron population advances at a median 2.215×10^8 neuron-steps/s, equivalent to 3.38 times real time. A routing-density sweep at the same population spans 100 to 10,000 contacts while holding source activity, external input width, and telemetry constant. Its 10,000-contact endpoint remains 2.17 times real time.




Figure 6: Connectivity scaling for a fixed population of 65,536 neurons. Each marker is the median of three p90 rates, and vertical ranges span three independently realized graphs. The line connects measurements with source events replayed at morphology- and role-conditioned target rates, isolating routing density from changes in activity. The orange diamond validates the 4,000-contact point with threshold-derived closed-loop firing. Every timed step includes the threshold scan, delayed routing, branch accumulation, an independently mutable efficacy for every contact, recurrent state, and telemetry for 1% of neurons; external afferent width remains fixed.

Population scale, connectivity, and simulated speed form a three-way tradeoff. The million-neuron measurement reports connectivity at a fixed population scale. The 65,536-neuron measurements extend the evidence to thousands of recurrent contacts and separate closed-loop execution from routing-density scaling. Appendix A.7 defines both workloads and reports the complete sweep.

5.2 AxoBench traces from recurrent population contexts

Recurrent feedback produces a distinct input distribution because event timing, excitation and inhibition, morphology, and graph structure jointly change branch load and temporal correlation. To sample that distribution, a thresholded AxoSim network generates branch-resolved contact histories, which are then replayed through matching detailed teachers. The recurrent learned system supplies the inputs, and five excitatory morphologies together with one inhibitory morphology supply detailed voltage and spike labels.

The axobench-population release contains 50,000 native-millisecond traces from 100 independently generated contexts. Eighty contexts provide 40,000 training traces, and 20 disjoint contexts provide 10,000 validation traces. Both splits preserve an 80/20 excitatory–inhibitory ratio. Validation recomputes split membership and checks manifest hashes, identifiers, array shapes, finite values, labels, and generator provenance. Paired perturbations reside in the separate axobench-interventions dataset, preserving the base population distribution.

Table 4: Composition of axobench-population. Complete recurrent contexts define the train-validation split.

Split and role	Traces	Contexts	Mean firing rate (Hz)
Train, excitatory	32,000	80	5.62
Train, inhibitory	8,000	80	6.32
Validation, excitatory	8,000	20	5.71
Validation, inhibitory	2,000	20	6.29
Total	50,000	100	–

5.3 Trace fidelity and recurrent validity are distinct

A plausible global firing rate can coexist with incorrect morphology-conditioned rates, temporal correlations, or late activity drift. A paired comparison isolates the effect of recurrent context by training one model on population-context examples and another on the original single-neuron distribution. On three unseen 4,000-neuron connectomes, population-context training reduces morphology discrepancy by 51.1%, temporal discrepancy by 7.2%, and stationarity discrepancy by 55.1%. It satisfies all three criteria fixed before evaluation, while the ordinary-input baseline misses the morphology criterion.




Figure 7: Development-corpus population-context training on three unseen 4,000-neuron connectomes. Thin lines pair training distributions on the same connectome; the purple diamond line connects cross-connectome means. Morphology measures mean squared log error in morphology-conditioned firing rates, temporal discrepancy measures lag-correlation error, and stationarity measures early-late log-ratio error. Annotations report the mean change and the number of improving connectomes. Lower is better throughout.

Training on all 50,000 traces improves held-out Voltage SERA from 299.99 to 285.84 mV², Dynamics SERA from 501.67 to 483.41 mV²/ms², exact-spike F1 from 0.136 to 0.163, and Mean F1 from 0.331 to 0.340. Yet the connected model misses the morphology-conditioned criterion on all three unseen connectomes. Consequently, improved trace-level metrics do not imply recurrent validity. Connected measurements determine whether the gain transfers to recurrent deployment.

6 Inference-time adaptation

6.1 Structural and synaptic transfer

Five leave-one-morphology-out comparisons measure specialization of a small structural adapter to an excluded morphology. Each source uses the other four morphologies, whereas target adaptation and evaluation use disjoint stimulus schedules from the omitted morphology. Updating one 58-value row improves Mean F1 and Dynamics SERA on all five targets. Voltage SERA improves on three, while the other two change by only +0.27% and +0.37%. Full tuning of 38,257 parameters provides an upper reference for the attainable gain.




Figure 8: Adaptation to five separately excluded morphologies. Each line follows one adaptation alternative across the five held-out targets. The 58-value adapter improves spike and dynamics fidelity for every morphology; full tuning provides an upper reference.

Where morphology transfer targets an omitted class, synaptic transfer addresses a known perturbation. Two conditions of one Hay layer-5 pyramidal-cell teacher differ only by an excitatory-conductance scale of 0.925 versus 1.0, with all 1,278 physical input channels. In the 0.925-to-1.0 direction, contact adaptation reduces Voltage SERA by 4.76%, reduces Dynamics SERA by 0.31%, and raises Mean F1 by 0.044. The learned contact changes correlate 0.724 with the known conductance intervention. In the reverse direction, correlation rises to 0.834 but Voltage SERA worsens. In the forward direction, contact coordinates both track and predictively compensate for the controlled shift. The reverse direction demonstrates the insufficiency of parameter correlation alone for a biological interpretation.




Figure 9: Transfer between matched conductance conditions of one detailed teacher. The 0.925-to-1.0 direction improves all three fidelity metrics. The reverse direction follows the imposed channel pattern more closely while increasing held-out Voltage SERA.

The role of each adaptive coordinate becomes concrete in these transfer comparisons. Updating a morphology row applies one class-shared correction, whereas synaptic updates change individual contact efficacies. Predictive improvement after a controlled shift demonstrates their practical utility. Cellular identification requires consistent predictive success across interventions.

6.2 Specialization to recurrent input distributions

Population-context inputs differ from the ordinary AxoBench distribution even when the neuron target is unchanged. Four AxoSim models and the released Branch-ELM family are adapted for 30,000 recurrent-domain presentations using the same context partition and architecture-specific schedules fixed on development data. Performance is then measured on 20 untouched target contexts.

Every AxoSim model improves target Mean F1. Mamba Medium reaches 0.412 after adaptation, compared with 0.136 for the strongest Branch-ELM target result. The other AxoSim models reach 0.364–0.402 and preserve the same large event-recovery separation. Target-domain results appear in Table 5; the full metric set is given in Appendix A.8.

Table 5: Spike fidelity after specialization to recurrent population inputs. The improvement is relative to the same model before target-domain adaptation.

Model	Target Mean F1	Improvement
AxoSim-GRU Small	0.379	+0.066
AxoSim-GRU Medium	0.402	+0.073
AxoSim-Mamba Small	0.364	+0.064
AxoSim-Mamba Medium	0.412	+0.094
Branch-ELM M100	0.136	+0.055

Corpus training targets a generally deployable recurrent model and therefore requires connected validation. Recurrent-domain adaptation instead measures whether exposed coordinates improve behavior in a specified target domain. The same data source supports both analyses under these distinct evaluation contracts.

6.3 Adaptation across a temporal task shift

Spiking Heidelberg Digits (SHD) supplies an external temporal distribution shift [4]. Development data compare seven optimizer–learning-rate–duration schedules and every nonempty combination of behavior, morphology, and synaptic groups. Morphology-plus-synapse adaptation, a learning rate of 2×10^{-3} , and 40 updates are fixed on one source model and reused for five independently trained sources. Each run uses 160 adaptation examples, 80 development examples, and 800 final examples, all disjoint.

Across six independently trained source models, target accuracy rises from 64.50% to 70.96%. The mean gain is 6.46 ± 1.20 percentage points, where the uncertainty is the sample standard deviation across six source models, and every target comparison improves. The source classifier used to choose the adaptation schedule reaches 91.88% before the temporal shift, localizing the lower target value to the shifted task.




Figure 10: SHD accuracy before and after inference-time adaptation. One source model fixes the adaptive parameter groups and update schedule; five independently trained models reuse it. Target accuracy improves in all six comparisons.

Accuracy gains are non-monotonic in contact count. At 700 contacts, two independently trained source models improve from 65.13% to 72.13%, a gain of 7.00 ± 1.06 points, while their mean source validation accuracy is 92.56%. The complete contact-density analysis appears in Appendix A.9. The external task therefore provides replicated evidence of useful within-system adaptation. Comparison with separate test-time-adaptation algorithms remains future work.

6.4 Population size, connectivity, and update time

The runtime contract measures one complete 256-step adaptation update. Each replay includes fresh pinned-host input transfer, exact sparse contact aggregation, forward execution, loss, full backward propagation, gradient clipping, and one capturable-Adam update in BF16. Real-time adaptation requires all seven changing replays to finish within 256 ms.

At 2,000 retained contacts per neuron, all seven updates remain below the deadline for 20,000 neurons, with a slowest update of 252.17 ms. The next measured population, 21,000 neurons, exceeds the deadline. Doubling connectivity to 4,000 contacts reduces the real-time population to 10,000 neurons,

whose slowest update takes 250.50 ms. Behavior, morphology, and synaptic parameters are optimized together, and every retained contact has an independently trainable efficacy.




Figure 11: Complete full-BPTT update time for densely connected populations. Markers show the slowest of seven changing updates at each population size; filled markers below the dotted 256-ms boundary satisfy the real-time criterion, while open markers show the first measured misses. Each update includes fresh pinned-host input, exact sparse aggregation, forward execution, loss, full backward propagation, gradient clipping, and one Adam step with all adaptive parameter groups enabled.

Repeated fixed-shape updates therefore support full-sequence learning at the reported population and connectivity. Section 8 discusses sustained long-horizon online learning as the next research direction.

7 Relation to learned and mechanistic simulators

7.1 Learned and mechanistic neuron models

Temporal neural networks can approximate the millisecond input–output map of a morphologically detailed cortical neuron and amortize repeated numerical integration of the teacher [2]. Spieler et al. introduced the Expressive Leaky Memory neuron as a compact alternative in which dendritic integration and slow memory reproduce detailed-neuron responses with relatively few parameters [15]. AxoSim uses released Branch-ELM models as its principal learned reference and the Hay layer-5 pyramidal-cell family as a principal mechanistic teacher [7].

Arbor and CoreNEURON instead accelerate networks of morphologically detailed neurons while retaining compartments, membrane mechanisms, and their mechanistic interpretation [1, 9]. Jaxley makes detailed simulation differentiable with respect to explicit morphological, synaptic, and ion-channel parameters [5]. These systems remain appropriate for studying the effect of a named conductance or compartment. AxoSim amortizes detailed teacher simulations by fitting each learned model once to generated responses and then reusing it for population realizations and gradient-based specialization within the learned domain.

7.2 Simulation abstractions and population execution

Brian 2 separates model equations from generated execution code, allowing a scientific description to target efficient backends [16]. BrainPy similarly combines dynamical systems, transformations, simulation, training, and analysis in one programming model [20]. The same separation is important for learned neurons. AxoSim exposes ordinary PyTorch modules for model construction and differentiation, then maps the compact population model to fused kernels while preserving contact identity, delay semantics, mutable efficacy, and parameter ownership.

Population execution also changes what must be measured. Cortical synapse counts vary by species, region, layer, and cell type, with relevant values in the thousands across reported preparations [14]. Accordingly, AxoSim attaches each population throughput measurement to a specific number of incoming synaptic contacts. Connected validation then measures morphology-conditioned responses, temporal correlations, and stationarity as a distinct level of recurrent evidence.

7.3 Metrics for learned neuron simulation

AxoBench introduced the core evaluation set used here, consisting of tolerance-dependent spike F1, Voltage SERA, and Dynamics SERA [21]. SERA integrates conditional squared error over teacher relevance and follows the imbalanced-regression construction of Ribeiro and Moniz [13]. Spike F1 is reported across explicit timing tolerances, complementing established spike-train distances [18, 19].

Each quantity measures a different aspect of the response. Spike F1 measures event recovery, Voltage SERA measures errors in relevant parts of the somatic trajectory, and Dynamics SERA measures errors in native-time voltage changes. A model can therefore receive different rankings on the three metrics. For recurrent use, morphology-conditioned activity, temporal correlation, and stationarity diagnostics extend the metric set to population scale.

8 Discussion and future work

8.1 Recurrent use changes the unit of validation

Learned simulation creates a hierarchy of validation levels. Held-out traces measure voltage and spike prediction under externally specified input. Connected evaluation then measures behavior after those predictions alter the next input distribution, while adaptation measures gradients through the same temporal feedback. Each level therefore requires its own evidence.

The distinction is visible in the population-data comparison. A controlled context corpus improves morphology, temporal, and stationarity behavior on unseen connectomes. Training on all 50,000 traces improves every held-out trace metric under the fixed training schedule, yet the connected model misses the morphology criterion. Small conditional errors can move recurrent state into a different input regime, and the graph then determines how those errors propagate. Consequently, a learned simulator should be validated at the level where its predictions become interventions.

These observations motivate an iterative division of labor. Detailed models generate teacher responses and controlled perturbations. Learned models amortize those responses across many connectomes or optimization steps. Connected diagnostics then identify distribution shifts and trigger new teacher simulations before the next deployment regime.

8.2 Edge-associated state dominates population scale

Parameter count of the shared neuron model is an incomplete predictor of population cost. At one million neurons, routing and branch accumulation grow from 59.7% to 78.8% of the component-median sum as incoming connectivity increases from 100 to 362 contacts. The fixed-activity routing sweep at 65,536 neurons then shows a broad plateau through 1,600 contacts, followed by a decline from 0.338 billion neuron-steps/s at 2,000 contacts to 0.142 billion at 10,000. Inference rate therefore depends jointly on population size, connectivity, activity, telemetry, and hardware.

Adaptation makes this relationship sharper because every retained edge also owns differentiable state. Doubling the retained contacts from 2,000 to 4,000 halves the measured real-time full-BPTT population from 20,000 to 10,000 neurons. Compact shared weights can be reused across the population, whereas edge-associated routing, storage, and gradients scale with the realized graph. Reports of learned population simulators should therefore attach neuron count to contact count and update mode.

8.3 Learned coordinates require predictive interventions

The adaptive parameter groups are attributable because their ownership is explicit. Morphology rows improve transfer to five omitted structures, and synaptic values can follow a known conductance pattern. Biological identification also requires predictive intervention. In the reverse conductance transfer, parameter correlation is higher even though Voltage SERA worsens. A learned coordinate gains biological relevance only when its change also predicts the response to a controlled intervention.

The intervention manifold supplies a complementary internal diagnostic. The association between manifold geodesics and detailed-teacher response distance rises from 0.211 in the branch representation to 0.446 after temporal integration. Because the input perturbation itself reaches 0.471, the learned geometry reflects both the imposed causal structure and its transformation by the temporal model. External response distances and predictive transfer together provide a stronger basis for interpretation than visual clustering alone.

8.4 Future work

Population-aware hyperparameter and objective selection on disjoint development connectomes should precede one evaluation on untouched connectomes. Such a comparison would determine whether trace-level gains transfer to recurrent behavior and whether mixed-domain training preserves broad single-neuron fidelity.

The AxoBench metrics are also unsaturated. Replicated training should examine depth, width, residual conditioning, objective balance, and compute while improving spike Mean F1, Voltage SERA, and Dynamics SERA together. In the measured frontier, the largest temporal core approaches the spike leader while smaller models lead the remaining metrics. Controlled, replicated trajectories can determine how depth, width, and compute produce that tradeoff.

Adaptation should extend from bounded transfer episodes to sustained long-horizon online learning. Such a study must measure cumulative task performance, memory growth, numerical stability, and drift over many updates. A matched-simulator comparison should fix morphology, stimulus, connectivity, precision, and output requirements across AxoSim and a detailed engine. Together, these analyses would strengthen both the scientific validity boundary and the practical interpretation of the runtime frontier.

9 Conclusion

AxoSim combines learned single-neuron fidelity with causal sparse population execution and differentiation through neuron- and contact-associated state. These capabilities are usually evaluated separately. Its strongest accuracy result is spike timing, accompanied by lower Voltage and Dynamics SERA than released Branch-ELM references. At population scale, a complete million-neuron step runs in real time at 362 contacts per neuron. A 65,536-neuron population remains 3.38 times faster than real time at 4,000 contacts, while complete full-BPTT updates remain real time for 20,000 neurons with 2,000 contacts per neuron.

The results also yield a methodological conclusion. Held-out trace gains and connected-population validity form two distinct levels of evidence because recurrent predictions change their own future inputs. AxoSim makes that distinction measurable by preserving contact semantics, exposing the ownership of adaptive state, and attaching each result to its measurement conditions.

Code and documentation. Installation, model and population interfaces, inference-time adaptation, worked examples, and the API reference are available at <https://axo.axym.org>.

References

- [1] Nora Abi Akar, Ben Cumming, Vasileios Karakasis, Anne Küsters, Wouter Klijn, Alexander Peyser, and Stuart Yates. Arbor—a morphologically-detailed neural network simulation library for contemporary high-performance computing architectures. *arXiv preprint arXiv:1901.07454*, 2019. URL <https://arxiv.org/abs/1901.07454>.
- [2] David Beniaguev, Idan Segev, and Michael London. Single cortical neurons as deep artificial neural networks. *Neuron*, 109(17):2727–2739.e3, 2021. doi: 10.1016/j.neuron.2021.07.002.
- [3] Mark M. Churchland, John P. Cunningham, Matthew T. Kaufman, Justin D. Foster, Paul Nuyujukian, Stephen I. Ryu, and Krishna V. Shenoy. Neural population dynamics during reaching. *Nature*, 487(7405):51–56, 2012. doi: 10.1038/nature11129.
- [4] Benjamin Cramer, Yannik Stradmann, Johannes Schemmel, and Friedemann Zenke. The heidelberg spiking data sets for the systematic evaluation of spiking neural networks. *IEEE Transactions on Neural Networks and Learning Systems*, 33(7):2744–2757, 2022. doi: 10.1109/TNNLS.2020.3044364.
- [5] Michael Deistler, Kyra L. Kadhim, Matthijs Pals, Jonas Beck, Ziwei Huang, Manuel Gloeckler, Janne K. Lappalainen, Cornelius Schröder, Philipp Berens, Pedro J. Gonçalves, and Jakob H. Macke. Jaxley: Differentiable simulation enables large-scale training of detailed biophysical models of neural dynamics. *Nature Methods*, 22(12):2649–2657, 2025. doi: 10.1038/s41592-025-02895-w.
- [6] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023. URL <https://arxiv.org/abs/2312.00752>.
- [7] Etay Hay, Sean Hill, Felix Schürmann, Henry Markram, and Idan Segev. Models of neocortical layer 5b pyramidal cells capturing a wide range of dendritic and perisomatic active properties. *PLOS Computational Biology*, 7(7):e1002107, 2011. doi: 10.1371/journal.pcbi.1002107.

- [8] Nikolaus Kriegeskorte, Marieke Mur, and Peter A. Bandettini. Representational similarity analysis—connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2:4, 2008. doi: 10.3389/neuro.06.004.2008.
- [9] Pramod Kumbhar, Michael Hines, Jeremy Fouriaux, Aleksandr Ovcharenko, James King, Fabien Delalandre, and Felix Schürmann. CoreNEURON: An optimized compute engine for the NEURON simulator. *Frontiers in Neuroinformatics*, 13, 2019. doi: 10.3389/fninf.2019.00063.
- [10] Leland McInnes, John Healy, and James Melville. UMAP: Uniform manifold approximation and projection for dimension reduction. *Journal of Open Source Software*, 3(29):861, 2018. doi: 10.21105/joss.00861.
- [11] Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. On the difficulty of training recurrent neural networks. In *Proceedings of the 30th International Conference on Machine Learning*, volume 28, pages 1310–1318, 2013. URL <https://proceedings.mlr.press/v28/pascanu13.html>.
- [12] Michael Pearce, Elana Simon, Michael Byun, and Daniel Balsam. Finding the tree of life in Evo 2. Goodfire Research, 2025. URL <https://www.goodfire.com/research/phylogeny-manifold>. Published August 27, 2025.
- [13] Rita P. Ribeiro and Nuno Moniz. Imbalanced regression and extreme value prediction. *Machine Learning*, 109(9–10):1803–1835, 2020. doi: 10.1007/s10994-020-05900-9.
- [14] Almut Schüz and Günther Palm. Density of neurons and synapses in the cerebral cortex of the mouse. *Journal of Comparative Neurology*, 286(4):442–455, 1989. doi: 10.1002/cne.902860404.
- [15] Aaron Spieler, Nasim Rahaman, Georg Martius, Bernhard Schölkopf, and Anna Levina. The expressive leaky memory neuron: An efficient and expressive phenomenological neuron model can solve long-horizon tasks. *arXiv preprint arXiv:2306.16922*, 2023. URL <https://arxiv.org/abs/2306.16922>.
- [16] Marcel Stimberg, Romain Brette, and Dan F. M. Goodman. Brian 2, an intuitive and efficient neural simulator. *eLife*, 8:e47314, 2019. doi: 10.7554/eLife.47314.
- [17] David Sussillo and Omri Barak. Opening the black box: Low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Computation*, 25(3):626–649, 2013. doi: 10.1162/NECO_a_00409.
- [18] M. C. W. van Rossum. A novel spike distance. *Neural Computation*, 13(4):751–763, 2001. doi: 10.1162/089976601300014321.
- [19] Jonathan D. Victor and Keith P. Purpura. Metric-space analysis of spike trains: Theory, algorithms and application. *Network: Computation in Neural Systems*, 8(2):127–164, 1997. doi: 10.1088/0954-898X/8/2/003.
- [20] Chaoming Wang, Tianqiu Zhang, Xiaoyu Chen, Sichao He, Shangyang Li, and Si Wu. BrainPy, a flexible, integrative, efficient, and extensible framework for general-purpose brain dynamics programming. *eLife*, 12:e86365, 2023. doi: 10.7554/eLife.86365.
- [21] Davide Wiest. AxoBench: Evaluating learned single-neuron surrogates as mechanistic simulators. Manuscript in preparation; public benchmark repository, 2026. URL <https://github.com/Axym-Labs/axobench>.

A Measurement contracts and supplementary evidence

A.1 Evaluation partitions

Table 6 gives the disjoint fitting, development, and final-evaluation units used by the principal comparisons.

Table 6: Data used for parameter fitting, schedule selection, and final evaluation.

Comparison	Fitting or adaptation	Development	Final evaluation
Ordinary AxoBench	training corpus	120 balanced traces	next 120: 24 threshold, 96 scored
Population specialization	64 recurrent contexts	16 recurrent contexts	20 untouched contexts
SHD adaptation	160 examples per source	80 examples on one source	800 examples on six sources
Connected populations	training traces	development connectomes	three paired unseen connectomes

A.2 Runtime measurement contracts

Table 7: Work included in each reported runtime quantity.

Quantity	Horizon and statistic	Included work	Excluded work
Isolated inference	500 steps, repeated median	model forward	data loading
Forward and backward	batch 32, 500 steps, 20 repeats	zero-grad, forward, loss, backward	loader, optimizer
Population inference	P4 advances, end-to-end p90	firing, routing, state, efficacy, 1% telemetry	data generation, initialization
Population adaptation	256 steps, seven replays	transfer, sparse input, forward, full BPTT, clip, Adam	graph capture

A.3 Model dimensions and optimization trajectories

The four full models use the dimensions below, which were fixed before the final training run.

Table 8: Temporal-core dimensions of the retained full models.

Model	Model width	State width	Layers	Head width
GRU Small	32	8	2	32
GRU Medium	128	32	2	128
Mamba Small	32	8	2	32
Mamba Medium	176	44	4	88

Training uses AdamW with batch size 32, initial learning rate 6×10^{-4} , cosine decay, and weight decay 10^{-4} . The objective combines scaled somatic-voltage MSE with weight 0.75, spike binary cross-entropy with weight 0.25, a direct SERA surrogate with weight 0.01, a differentiable F1 tolerance-curve term over 0–5 steps with weight 0.05, and near-spike voltage error within five steps with weight 0.002. The SERA surrogate evaluates 101 relevance thresholds.

Trained states at 24, 32, and 40 million presentations are compared on 120 development traces. GRU Small is retained at 32 million presentations, whereas the other three profiles are retained at 40 million. A separate set of 120 morphology-balanced traces supplies the final evaluation block.




Figure 12: Paired inference and forward-plus-backward throughput for one representative of each architecture family. Each rate is reconstructed from the stored median latency, batch size, and 500-step window; loader and optimizer work are excluded.




Figure 13: Interval mean optimization objective for the four stable models after the common training run resumed near 19.5 million presentations. The interval view distinguishes late plateauing and batch variation from numerical divergence.

A.4 Intervention-response geometry

The intervention analysis uses the final AxoSim–Mamba Medium model and all 12,500 paired traces in axobench–interventions. The five intervention families are event dropout, excitatory dropout, inhibitory dropout, site silencing, and temporal jitter, crossed with five morphologies and eight dataset shards per combination. Baseline state is subtracted from intervention state for every pair. The first 500 ms are excluded, and the remaining 5,500 ms are averaged into eleven consecutive windows.

The signal column in Figure 3 uses the 500-ms window with the highest teacher spike count among one deterministically selected pair from each of the 200 shards. This trace exposes measured input, state, and output signals in the architecture diagram. It is excluded from every reported aggregate statistic. The input display uses 10-ms event counts, while the temporal representation, voltage, and spike probability are sampled at the native 1-ms resolution. The displayed probability is the sigmoid of each individual logit. PC 1 is computed after standardizing each temporal feature within this window.

Panel (a) groups each shard deterministically into disjoint mini-cohorts of seven or eight pairs, producing 1,575 points. Each point concatenates the mean 176-dimensional temporal-state displacement in the eleven windows. Features are standardized before fitting a three-dimensional cosine-distance UMAP with 30 neighbors, minimum distance 0.12, and random seed 20260930. UMAP coordinates are used only for the display.

Panel (b) selects one pair from each shard by hashing the seed, intervention, and shard name. State displacements are averaged into 500-ms windows, and one PCA is fitted across every selected trace and time. The first three components explain 53.2%, 21.1%, and 8.9% of the variance. Condition paths average 40 selected traces. For the gray flow field, each individual transition is assigned to a cell in a $4 \times 4 \times 4$ grid spanning the 3rd–97th percentile range. Arrows report mean displacement in occupied cells with at least eight transitions; the 32 most populated cells are shown.

The external-ruler comparison averages the eight shards within each of the 25 groups defined by intervention and morphology. In each representation, feature-standardized angular distances form the smallest connected symmetric nearest-neighbor graph. The required neighbor counts are five for branch state, six for temporal state, and five for the input descriptor. Shortest-path distances are then compared by Spearman correlation with Euclidean distance between standardized detailed-teacher signatures. Each teacher signature contains eleven mean voltage changes and eleven spike-count changes. The 95% intervals come from 1,000 resamples of the eight shards within every group; each resample reconstructs the group representations and graph. The direction-consistency margin subtracts the mean cosine similarity across interventions within a morphology from the mean similarity across morphologies within an intervention.

A.5 Central figure values and axis mappings

Table 9 gives the values represented in Figure 1. Inference throughput uses the absolute mapping $1 - \exp(-r/10^7)$, with r in neuron-steps/s. Voltage and Dynamics SERA use $\exp(-E/350)$ and $\exp(-E_\Delta/450)$, respectively, so zero error maps to one and unbounded error approaches zero. Mean F1 and F1 at 5 ms are untransformed. These mappings are fixed independently of the compared values.

CoreNEURON supplies the target trajectories, so its self-comparison has zero SERA and unit F1. The simulation rate comes from a different hardware and workload record and appears as an absolute measurement rather than the denominator of a cross-hardware speedup. Learned-model rates use the shared 500-step inference contract described in Section 3.

Table 9: Raw values used by the central figure. Rates are millions of neuron-steps/s; Voltage SERA is in mV^2 , and Dynamics SERA is in mV^2/ms^2 .

Variant	Inference	Voltage SERA	Dynamics SERA	Mean F1	F1 at 5 ms
CoreNEURON reference	0.00085	0	0	1.000	1.000
Branch-ELM M30	0.742	307.41	448.57	0.077	0.102
AxoSim-Mamba Medium	4.396	244.46	439.00	0.520	0.641
AxoSim-GRU Small	21.786	243.92	429.12	0.486	0.604

The learned-model comparison concerns fitted quality and measured execution rates. AxoSim and the released Branch-ELM models received different upstream optimization budgets, so these measurements do not establish relative training efficiency.

A.6 Pareto-arrow selection and FLOP contract

The orange arrow in each panel of Figure 2 follows one fixed rule. Within the Branch-ELM group and the combined AxoSim group, points dominated on both the displayed efficiency and fidelity axes are discarded. The remaining coordinates are min-max normalized within their group after orienting both axes so that larger values are better. The selected point is the nondominated candidate with the smallest Euclidean distance to the ideal corner $((1,1))$, with model identifier as a deterministic tie-breaker. The precomputed endpoint table in `data/tikz/pareto-arrows.csv` decouples selection from figure layout.

Forward-pass FLOPs count the operations recognized by `torch.profiler` for one batch-one, 500-step sequence with 1,278 input features. The same measurement is applied to every learned model. Parameter counts include trainable parameters only. Neither quantity includes data loading or optimizer operations.

A.7 Population timing and connectivity

Both population measurements use one RTX 5090 and BF16 compact-neuron execution at 1 ms resolution. Each timed step includes delayed routing that preserves source identity, mutable positive efficacies, recurrent state, and pinned-host voltage telemetry for 1% of neurons. The million-neuron measurement uses closed-loop spike discovery and 362 retained contacts per neuron. Three target-activity replicates span 1.174–1.197 billion neuron-steps/s, with a median of 1.177 billion neuron-steps/s.

Figure 6 uses 65,536 neurons to measure denser graphs at 100, 200, 400, 800, 1,600, 2,000, 3,200, 4,000, 6,400, 8,000, and 10,000 retained recurrent contacts per neuron. A threshold scan remains in the timed step, after which source events are replayed at fixed morphology- and role-conditioned target rates. This replay holds activity constant while the measurement varies routing density. External input width remains fixed at 1,278 afferents, and every recurrent contact retains an independently mutable efficacy. Repeated source-target pairs are permitted, as multiple synapses may connect the same neuron pair.

Each contact count is measured on three independently realized graphs. The reported value is the median p90 step rate, with the full replicate range shown by the error bars. In this fixed-activity routing sweep, the 4,000-contact median is 249.87 million neuron-steps/s and the 10,000-contact median is 142.09 million neuron-steps/s, equivalent to 2.17 times real time. The curve connects measured medians and terminates at the 10,000-contact endpoint.

The orange 4,000-contact marker replaces source replay with threshold-derived closed-loop firing. Across three graph realizations, the measured-to-target rate ratio spans 1.021–1.026 with fixed thresholds throughout the measured window. The median p90 rate is 221.47 million neuron-steps/s, equivalent to 3.38 times real time for 65,536 neurons.

Quantitative mouse-neocortex estimates are on the order of thousands of contacts per neuron, with a reported average near 8,200 and substantial region and cell-type variation. Contacts also differ from counts of unique presynaptic partners. Population size, connectivity, and simulated time per wall clock second should therefore be reported together.

A 426-contact implementation used a stride-three bit-mask selector. Bit-mask reduction is valid only for power-of-two moduli; with a non-power-of-two fanout it collapses residue classes and does not realize the declared destinations. Correct modulo routing restores every class and yields more scattered target addresses. In the directly comparable records, median routing time rises from 1.579 to 1.965 ms. The increase is consistent with both the arithmetic remainder and the less-local target addresses; the timing record does not identify their individual contributions.

A.8 Population-input specialization details

Table 10: Target-domain metrics after 30,000 population-context presentations. SERA values use the same squared units as the ordinary AxoBench metrics.

Model	Mean F1	Voltage SERA	Dynamics SERA
AxoSim-GRU Small	0.379	48.10	20.24
AxoSim-GRU Medium	0.402	50.34	23.49
AxoSim-Mamba Small	0.364	47.89	21.34
AxoSim-Mamba Medium	0.412	49.23	24.23
Branch-ELM M50	0.132	36.87	16.36
Branch-ELM M100	0.136	36.93	16.36

Recurrent-domain specialization complements the ordinary-domain comparison. Section 4 gives the direct cross-family SERA result on ordinary AxoBench. The target-domain analysis instead measures how each model adapts its spike behavior to recurrently generated inputs.

A.9 SHD contact-density analysis




Figure 14: SHD adaptation as the number of synaptic contacts changes. The development result on one source is non-monotonic; two independently trained sources confirm the selected 700-contact setting.

A.10 Connected diagnostic definitions

Let $r_{s,m}$ denote the firing rate of excitatory morphology group m over either the full or tail window s , and let r_m^* denote its contract target. The morphology discrepancy is

$$D_{\text{morph}} = \frac{1}{2|\mathcal{M}|} \sum_{s \in \{\text{full}, \text{tail}\}} \sum_{m \in \mathcal{M}} \left[\log \left(\frac{r_{s,m}}{r_m^*} \right) \right]^2.$$

For per-neuron input-count lag correlations $\rho_{s,\ell}$, evaluated at $\ell \in \{1, 4, 16, 64, 256\}$ ms, the temporal discrepancy is

$$D_{\text{temp}} = \frac{1}{2|\mathcal{L}|} \sum_s \sum_{\ell \in \mathcal{L}} \left(\frac{\rho_{s,\ell} - \rho_\ell^*}{0.2} \right)^2.$$

Finally, let q range over global rate, excitatory rate, inhibitory rate, and input-count coefficient of variation. The stationarity discrepancy is

$$D_{\text{stat}} = \frac{1}{4} \sum_q \left[\log \left(\frac{q_{\text{tail}}}{q_{\text{full}}} \right) \right]^2.$$

Targets and normalizations are fixed before held-out testing. The three quantities remain separate throughout evaluation.

A.11 Conditional single-neuron examples




Figure 15: Balanced conditional examples from the final comparison block. The upper row is the spike-containing 300 ms window with the largest voltage-MSE advantage for the selected AxoSim-Mamba model over the strongest Branch-ELM spike baseline; the lower row is the largest disadvantage under the same rule. These selected examples expose one success and one failure and are not estimates of average performance. Black denotes the detailed teacher, solid purple AxoSim, and dashed cyan Branch-ELM.

A.12 Inference and full-BPTT contact memory




Figure 16: Peak CUDA memory per neuron for inference and full BPTT across dense branch expansion, direct einsum, and direct BMM contact-input implementations at 24 contacts and 256 timesteps.

Dense branch expansion uses approximately 2.686 MB per neuron for inference and 2.713 MB for full BPTT. Direct einsum and direct BMM each use approximately 0.201 MB for inference and 0.326 MB

for full BPTT. Normalized per-neuron memory permits comparison across the different executable population sizes; the source CSV also records raw populations and byte counts.

A.13 Adaptation timing boundary

The dense-contact benchmark optimizes behavior, morphology, and synaptic parameters together through 256 native-millisecond steps. At 2,000 retained contacts per neuron, 20,000 neurons pass all seven changing updates and 21,000 exceed the 256-ms deadline. At 4,000 contacts, 10,000 neurons pass and 11,000 exceed the deadline. The reported curve uses the slowest update at each population size. Timing includes fresh pinned-host input, exact sparse contact aggregation, transfer, forward execution, loss, full backward propagation, gradient clipping, and Adam; one-time graph capture is excluded.